



# SPIKING NEURAL NETWORKS

Daniel Price and Brandon Nguyen

# Artificial Neural Networks

- Poorly represent the human brain
- Matrix Multiplication (GEMM)
  - More Computation -> More Energy
- Offline learning
- Backpropagation
  - High training Cost

# Spiking Neural Networks

- Incorporates neuron models
- Spikes convert GEMM into accumulation
  - Less Computation -> Less Energy
- Online learning
- Surrogate Gradient Descent / STDP (Unsupervised)
  - Low training cost

# Applications

- **Robotics**
- Neuroscience
- **Autonomous systems**
- Natural language processing
- **Computer vision**
- etc.



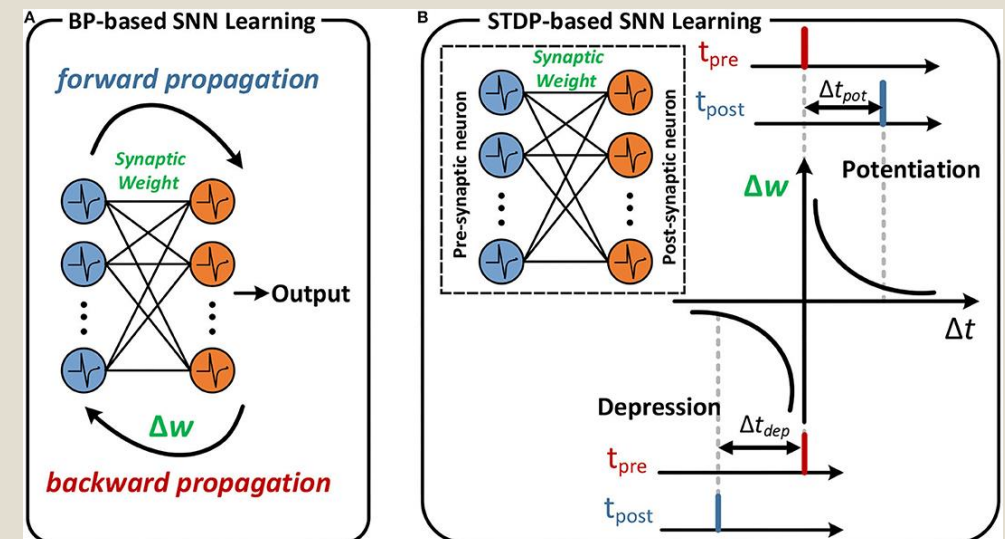
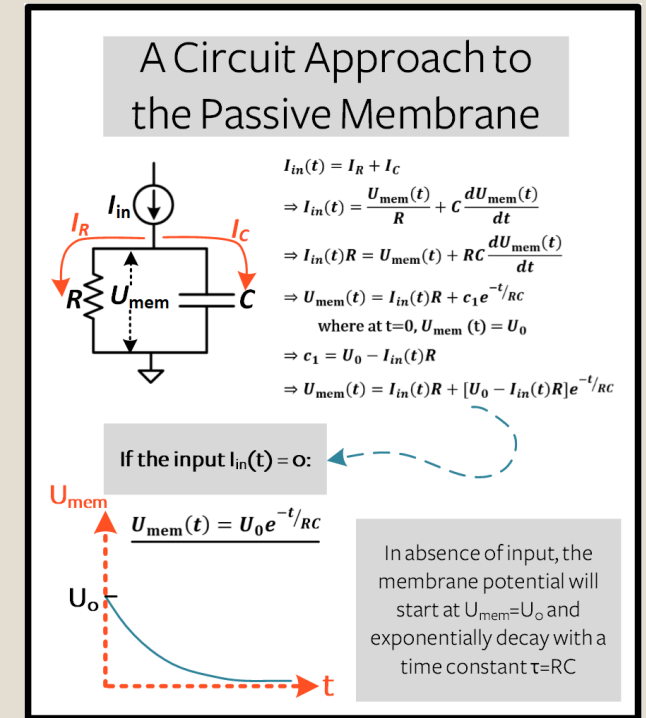
Canada's Mars Exploration Science Rover

# Leaky-Integrate and Fire

- Simple Neuron Model
- Represents neurons as an RC circuit
- Input current increase membrane potential
- When membrane potential breaks threshold, potential resets

## Learning

- Surrogate Gradient Descent
  - Step function learning (spike or no spike)
  - Replace derivative with a surrogate function
- STDP (Spike-Time Dependent Plasticity)
  - Unsupervised learning
  - Updates weights on spiking behavior
    - Presynaptic / Postsynaptic spikes



# Workflow Distribution

## Brandon:

- Visualization of model architecture
- Data Processing and Data Visualization construction
- Code Documentation

## Daniel:

- CNN and CSNN Architecture Design
- Generation of results and case study data
- Hyperparameter tuning

# Problem Description

Train a SNN on MNIST  
Dataset for a supervised  
multi-class classification task.



## Evaluation Goals:

Compare with ANN /  
CNN on **accuracy** and  
**training loss**

Compare computation  
cost in **FLOPs** (Floating  
Point Operations)

# Initial Approach

- LIF / STDP model
- Struggled to correlate images to labels

## Why?

- STDP struggles with classification
  - Unsupervised learning
  - Learning does not use labels
- STDP is used more commonly in TNNs (Temporal Neural Networks)
- SNNs are emerging:
  - Limited application
  - Research focused

# Attempts to salvage Initial Approach


## Normalization & Regularization

- Varying threshold voltage
- Reward modulated STDP (R-STDP)
  - Rewards/Punishes neurons based on prediction

## Hyperparameter tuning variation

- R/C/Tau
- Timesteps
- Learning Rate

## Batch Normalization (BN)



- $\mu$ : mean of  $x$  in mini-batch
- $\sigma$ : std of  $x$  in mini-batch
- $\gamma$ : scale
- $\beta$ : shift
- $\mu, \sigma$ : functions of  $x$ , analogous to responses
- $\gamma, \beta$ : parameters to be learned, analogous to weights

**Image:** <http://tzutalin.blogspot.com/2017/07/deep-learning-batch-normalization-note.html>

# New Approach

- **BNTT (Batch Normalization Through Time)**
  - Codebase detailing implementation
  - Batch Normalization
    - Normalizes inputs across batch
    - Reduces covariant shift
    - Applies to each timestep
  - Backpropagation
    - Preferable with SNNs / Classification
- **Solid foundation, but more complex than necessary.**

**Code:** <https://github.com/Intelligent-Computing-Lab-Yale/BNTT-Batch-Normalization-Through-Time/blob/main/model.py>



## New Approach: Our Implementation

### Simplify Architecture:


- Reduce layers from 9 to 3
- Reduce Shape / Input channels
  - Match input (MNIST)
  - Lower Complexity (hidden shapes)
- Reduce epochs and timesteps

### Justification:

- Model is too complex for dataset

# Experimental Methodology

- Dataset: MNIST Handwritten Digits
- Model Architecture
- Hyperparameters used
- Encoding type
- Case Study
- Evaluation: Classification accuracy, training loss, and FLOPs



MNIST Examples

# MNIST Dataset

A database of handwritten digits containing a training set of 60,000 examples, and a test set 10,000 examples.



Using 10% of the dataset due to training overhead (SNN)

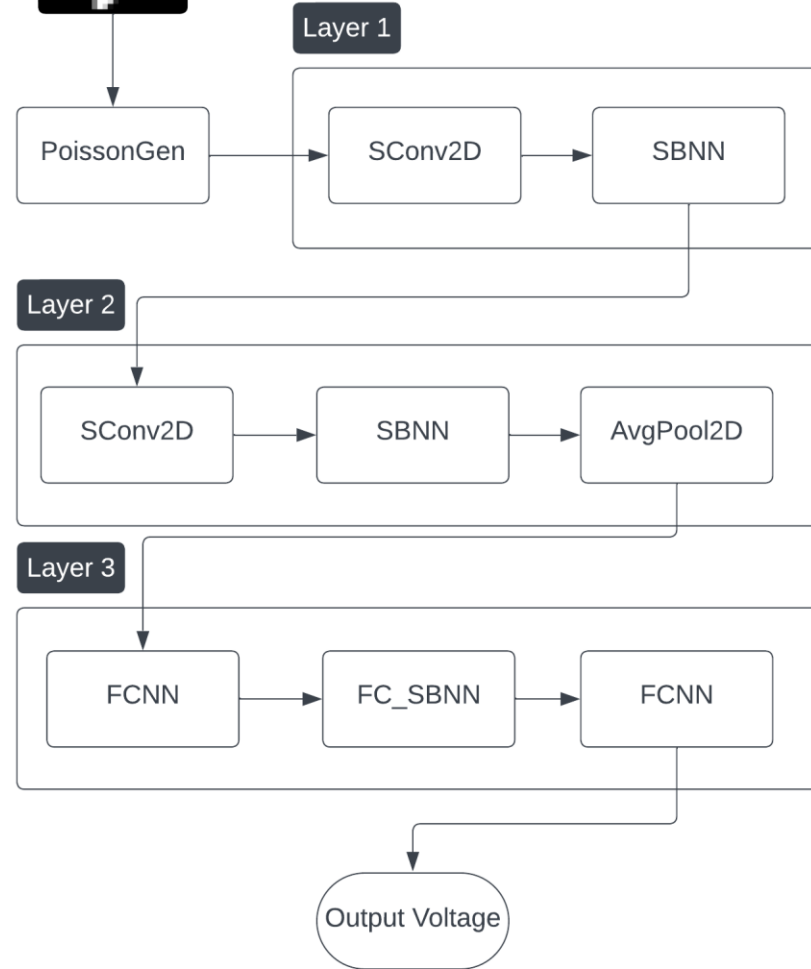
Training on 6,000  
samples.

Testing on 1,000  
samples.



**Goal:** Correctly classify images as one of the ten digits.

Input: MNIST Batch



# Model

- Visual diagram of our model architecture and flow of our forward pass function if observing layers only.
- Baseline CNN ignores PoissonGen and uses non-spiking equivalents
  - SConv2D -> Conv2D

# Hyperparameters - Baseline

- **Global Hyperparameters:**
  - batch size, learning rate, number of epochs
- **Convolutional Layer 1**
  - Input Channels: 1; Output Channels 64; Kernel Size: 1; Stride: 1
- **Convolutional Layer 2**
  - Input Channels: 64; Output Channels: 64; Kernel Size: 3; Stride: 1
- **2D Batch Normalization between Layers**
  - Input Channels: 64; Epsilon: 1e-4; Momentum: 0.1;
- **Pooling (after L2 and BN)**
  - Average Pooling; Kernel Size: 2
- **Fully-connected Layer 1**
  - Input Dimension: 14 \* 14 \* 64; Output Dimension: 512
- **1D Batch Normalization**
  - Input Dimension: 512; Epsilon: 1e-4; Momentum: 0.1
- **Fully-connected Layer 2**
  - Input Dimension: 512; Output Dimension: 10
- **Activation Function**
  - Rectified Linear Unit following each BN operation.

# Hyperparameters - CSNN

## **Global Hyperparameters:**

- number of time steps, leakage memory, batch size, learning rate, number of epochs

## **Convolutional Layer 1**

- Input Channels: 1; Output Channels 64; Kernel Size: 1; Stride: 1

## **Convolutional Layer 2**

- Input Channels: 64; Output Channels: 64; Kernel Size: 3; Stride: 1

## **Spiking 2D Batch Normalization between Layers**

- Input Channels: 64; Epsilon: 1e-4; Momentum: 0.1;

## **Pooling (after L2 and BN)**

- Average Pooling; Kernel Size: 2

## **Fully-connected Layer 1**

- Input Dimension: 14 \* 14 \* 64; Output Dimension: 512

## **Spiking 1D Batch Normalization**

- Input Dimension: 512; Epsilon: 1e-4; Momentum: 0.1

## **Fully-connected Layer 2**

- Input Dimension: 512; Output Dimension: 10

## **Activation Function**

- N/A

# Encoding Type

## Rate Encoding

- Represents data as firing rate (frequency)

## Latency Encoding

- Represents data as spike timings

## BNTT uses rate encoding

- SNNs prefer Rate encoding
- TNNs prefer Latency encoding

# Case Study

## Hyperparameter Variation


- Batch Size
- Epochs
- Timesteps
  - Steps in each forward pass of a neuron
- Leakage Voltage
  - Percentage of voltage retained between each step

## Metrics


- Testing Accuracy
- Total FLOPs (Normalized)



# Case Study – Testing Accuracy vs. Batch Size




CNN Results:




CSNN Results:

# Case Study – Testing Accuracy vs Learning Rate




CNN Results:




CSNN Results:

# Case Study – Testing Accuracy vs Epochs





CNN Results:





CSNN Results:

# Case Study – Testing Accuracy vs Leakage Voltage / Time Steps



**CSNN Results:**

# Case Study – Norm FLOPs vs Leakage Voltage / Time Steps



CSNN Results:

Model	CNN	SNN
Batch Size	32	32
Learning Rate	0.1	0.2
Epochs	50	50
Time Steps	N/A	10
Leakage Voltage	N/A	0.985
Training Loss	0.00098	0.00371
Testing Accuracy	98.535	98.145
Normalized Inference FLOPs	1x	0.428x

## Evaluation

- Training Loss
- Classification Testing Accuracy
- Floating Point Operations comparison

## Takeaway

- CSNN Model classifies with similar accuracy
- Computes 42.8% flops compared to CNN
- Ignores multiplier, majority of computational energy
- Achieves Significant decrease in energy usage.

QUESTIONS?